

Shuqi Hu

shuqi.hu@sv.cmu.edu

(669)2643158

<https://www.linkedin.com/in/shuqi-hu>

3200 Zanker Road, San Jose

PROFILE

Masters student seeking full-time positions of Software Development Engineer, good at Distributed Systems and Machine Learning.

SKILLS

Programming Languages: Java / Python / C / Bash / Go

Platforms & Libraries: Hadoop / Solr / Lucene / ZooKeeper / MapReduce / Avro / Grafana / AWS / TensorFlow

Databases: MySQL / Oracle / Cassandra / Redis Cluster / Memcached / Hive / Prometheus / MongoDB

Software Developments: Git / Perforce / Maven / Gradle / Unit Test / Spring Boot

EDUCATION

Carnegie Mellon University, Mountain View, CA

Expected 12/2018

M.S. in Electrical and Computer Engineering (GPA: 3.52/4.0)

Tongji University, Shanghai, China

07/2017

B.S. in Electronics & Information Engineering (GPA: 3.72/4.0)

EXPERIENCE

Electronic Arts Inc., Redwood City, CA

05/2018 – 08/2018

Software Engineer Intern, EADP Data & AI

- Deployed a POC search engine which indexes player profiles and serves game teams with analytical insights using SolrCloud and ZooKeeper aimed to replace the current Lucene index system
- Developed a Solr Plugin to index Avro files and streams and a language parser to translate Boolean expressions to Solr queries using Antlr, kept the new system completely compatible with the current system
- Created an interface using Prometheus and Grafana to monitor and alert on the cluster's status
- Improved the current Lucene index system with fault tolerance, high scalability and faster search response, thus saving half the AWS EC2 machines used in the current system

Huawei Technologies Co., Ltd., China

07/2016 – 08/2016

Software Engineer Intern, Enterprise Cloud Business Department

- Participated in the daily maintenance work in the data center to resolve the failures in the server connection
- Deployed the client devices and secured connection with data center for Huawei Cloud Desktop Services
- Assisted the team to expand market share, received group recognition for system management and monitor

PROJECT

Restaurant Search and Recommendation Java Web Service

- Developed a dynamic web page and Android App for restaurants search and content-based recommendation
- Created Java servlets with RESTful APIs to handle HTTP requests, deployed server on AWS EC2 and built relational and NoSQL databases (MySQL, MongoDB) to capture real restaurant data from Yelp API
- Personalized search and recommendation using collaborative filtering based on view history and favorites

Distributed File System

- Implemented a write-intensive distributed photo storage solution separating name nodes and data nodes
- Built the name node cluster using Cassandra maintaining meta data of photos and data nodes and the data node cluster based on Linux file system with logical store to physical store mapping
- Combined cache using Redis Cluster in front of both clusters with 72% hit rate and 15x faster read requests

Hadoop Streaming to Analyze Nov 2016 Wikipedia

- Processed a large text dataset using parallel programming running in Hadoop frameworks on AWS
- Implemented a robust Java solution using the MapReduce programming model and remote debug method
- Analyzed the central tendency, dispersion and shape of the dataset's distribution, figured out the most popular monthly and daily articles and topics on Wikipedia for Nov 2016